



DataBricks Spark – Azure DataBricks

Contents

1. **Course Introduction**
2. **Why Apache Spark?**
3. **Spark Cluster Managers**
4. **Introduction to DataBricks**
5. **DataBricks Components**
6. **RDD**
7. **Transformations and Actions in RDD**
8. **DataFrame**
9. **Transformation and Actions in Dataframe**
10. **Working with DataFrames**
11. **SparkSQL**
12. **File systems and sources supported by Spark**
13. **DeltaLake**
14. Spark Applications
15. Batch ETL using Spark
16. Introduction to Kafka
17. Real-Time ETL and Event partition using Kafka and Spark
18. Spark MLLib and Machine Learning using Spark

Submit Options

- The Spark submit script provides many options to specify how the application should run
 - Most are the same as for pyspark2 and spark2-shell
- General submit flags include
 - master: local, yarn, or a Mesos or Spark Standalone cluster manager URI
 - jars: Additional JAR files (Scala and Java only)
 - pyfiles: Additional Python files (Python only)
 - driver-java-options: Parameters to pass to the driver JVM

- num-executors: Number of executors to start application with
- driver-cores: Number cores to allocate for the Spark driver
- queue: YARN queue to run in

- Spark determines how to partition data in an RDD, Dataset, or DataFrame when
 - The data source is read
 - An operation is performed on a DataFrame, Dataset, or RDD
 - Spark optimizes a query
 - You call repartition or coalesce

- Partitions are determined when files are read
 - Core Spark determines RDD partitioning based on location, number, and size of files
 - Usually each file is loaded into a single partition
 - Very large files are split across multiple partitions
 - Catalyst optimizer manages partitioning of RDDs that implement DataFrames and Datasets
- You can view the number of partitions in an RDD by calling the function `getNumPartitions`

```
myRDD.getNumPartitions
```

Language: Scala

```
myRDD.getNumPartitions()
```

Language: Python


Spark Stages and Tasks

- A task is a series of operations that work on the same partition and are pipelined together
- Stages group together tasks that can run in parallel on different partitions of the same RDD
- Jobs consist of all the stages that make up a query
- Catalyst optimizes partitions and stages when using DataFrames and Datasets
 - Core Spark provides limited optimizations when you work directly with RDDs
 - You need to code most RDD optimizations manually
 - To improve performance, be aware of how tasks and stages are executed when working with RDDs

Spark Execution Plan

- Spark creates an execution plan for each job in an application
- Catalyst creates SQL, Dataset, and DataFrame execution plans
 - Highly optimized
- Core Spark creates execution plans for RDDs
 - Based on RDD lineage
 - Limited optimization

Spark Execution Plan



Spark Execution Plan

- Spark constructs a DAG (Directed Acyclic Graph) based on RDD dependencies
- Narrow dependencies
 - Each partition in the child RDD depends on just one partition of the parent RDD
 - No shuffle required between executors
 - Can be pipelined into a single stage
 - Examples: map, filter, and union

Spark Execution Plan

- Wide (or shuffle) dependencies
 - Child partitions depend on multiple partitions in the parent RDD
 - Defines a new stage
 - Examples: reduceByKey, join, and groupByKey

Controlling the Number of partitions

- Partitioning determines how queries execute on a cluster
 - More partitions = more parallel tasks
 - Cluster will be under-utilized if there are too few partitions
 - But too many partitions will increase overhead without an offsetting increase in performance
- Catalyst controls partitioning for SQL, DataFrame, and Dataset queries
- You can control how many partitions are created for RDD queries

Controlling the Number of partitions

- Specify the number of partitions when data is read
 - Default partitioning is based on size and number of the files (minimum is two)
 - Specify a different minimum number when reading a file

```
myRDD = sc.textFile(myfile,5)
```

Controlling the Number of partitions

- Manually repartition
 - Create a new RDD with a specified number of partitions using repartition or coalesce
 - coalesce reduces the number of partitions without requiring a shuffle
 - repartition shuffles the data into more or fewer partitions

```
newRDD = myRDD.repartition(15)
```

Controlling the Number of partitions

- Specify the number of partitions created by transformations
 - Wide (shuffle) operations such as reduceByKey and join repartition data
 - By default, the number of partitions created is based on the number of partitions of the parent RDD(s)
 - Choose a different default by configuring the spark.default.parallelism property

```
spark.default.parallelism 15
```

Catalyst Optimizer

- Catalyst can improve SQL, DataFrame, and Dataset query performance by optimizing the DAG to
 - Minimize data transfer between executors
 - Such as broadcast joins—small data sets are pushed to the executors where the larger data sets reside
 - Minimize wide (shuffle) operations
 - Such as unioning two RDDs—grouping, sorting, and joining do not require shuffling
 - Pipeline as many operations into a single stage as possible
 - Generate code for a whole stage at run time
 - Break a query job into multiple jobs, executed in a series

Catalyst Optimizer

- Execution plans for DataFrame, Dataset, and SQL queries include the following phases
 - Parsed logical plan—calculated directly from the sequence of operations specified in the query
 - Analyzed logical plan—resolves relationships between data sources and columns
 - Optimized logical plan—applies rule-based optimizations
 - Physical plan—describes the actual sequence of operations
 - Code generation—generates bytecode to run on each node, based on a cost model


Spark persistence

- ❑ You can persist a DataFrame, Dataset, or RDD
 - Also called caching
 - Data is temporarily saved to memory and/or disk
- ❑ Persistence can improve performance and fault-tolerance
- ❑ Use persistence when
 - Query results will be used repeatedly
 - Executing the query again in case of failure would be very expensive
- ❑ Persisted data cannot be shared between applications

Spark persistence

```
over20DF = spark.read. \  
  json("people.json"). \  
  where("age > 20") \  
pcodesDF = spark.read. \  
  json("pcodes.json") \  
joinedDF = over20DF. \  
  join(pcodesDF, "pcode")
```

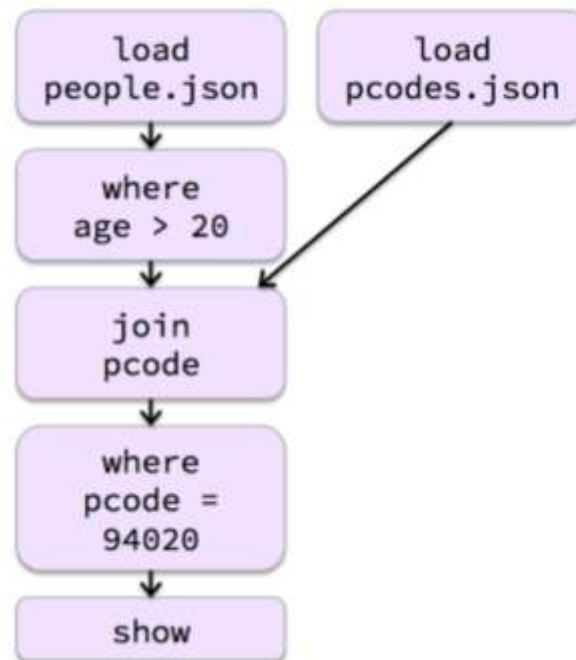
Language: Python



Spark persistence

```
over20DF = spark.read. \
    json("people.json"). \
    where("age > 20")
pcodesDF = spark.read. \
    json("pcodes.json")
joinedDF = over20DF. \
    join(pcodesDF, "pcode")
joinedDF. \
    where("pcode = 94020"). \
    show()
```


Language: Python



Spark persistence

```
over20DF = spark.read. \
    json("people.json"). \
    where("age > 20")
pcodesDF = spark.read. \
    json("pcodes.json")
joinedDF = over20DF. \
    join(pcodesDF, "pcode"). \
    persist()
```

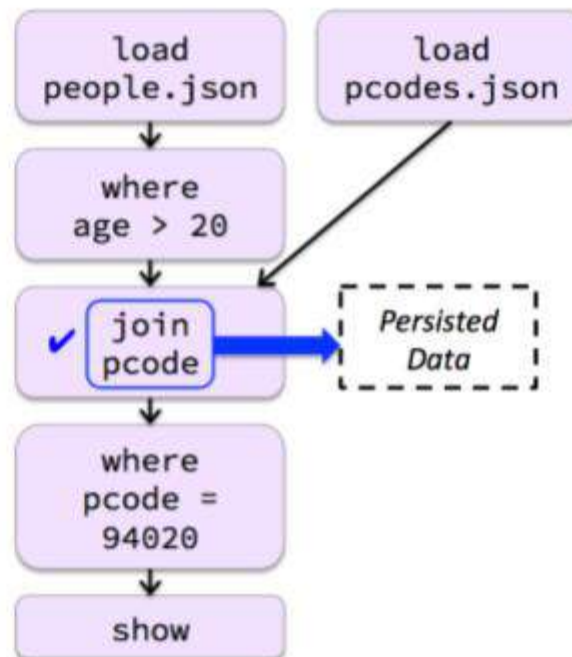
Language: Python



Spark persistence

```
over20DF = spark.read. \
    json("people.json"). \
    where("age > 20")
pcodesDF = spark.read. \
    json("pcodes.json")
joinedDF = over20DF. \
    join(pcodesDF, "pcode"). \
    persist()
joinedDF. \
    where("pcode = 94020"). \
    show()
```

Language: Python



Spark persistence

```
over20DF = spark.read. \
    json("people.json"). \
    where("age > 20")
pcodesDF = spark.read. \
    json("pcodes.json")
joinedDF = over20DF. \
    join(pcodesDF, "pcode"). \
    persist()
joinedDF. \
    where("pcode = 94020"). \
    show()
joinedDF. \
    where("pcode = 87501"). \
    show()
```

Language: Python

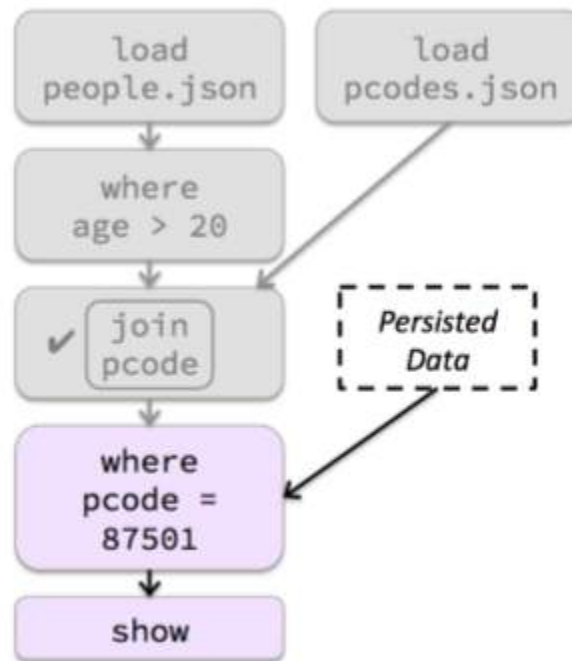


Table and View Persistence

- ❑ Tables and views can be persisted in memory using CACHE TABLE
`spark.sql("CACHE TABLE people")`
- ❑ CACHE TABLE can create a view based on a SQL query and cache it at the same time
`spark.sql("CACHE TABLE over_20 AS SELECT * FROM people WHERE age > 20")`
- ❑ Queries on cached tables work the same as on persisted DataFrames, Datasets, and RDDs
 - The first query caches the data
 - Subsequent queries use the cached data

Persistence Storage levels

- ❑ Storage levels provide several options to manage how data is persisted
 - Storage location (memory and/or disk)
 - Serialization of data in memory
 - Replication
- ❑ Specify storage level when persisting a DataFrame, Dataset, or RDD
 - Tables and views do not use storage levels
 - Always persisted in memory
- ❑ Data is persisted based on partitions of the underlying RDDs
 - Executors persist partitions in JVM memory or temporary local files
 - The application driver keeps track of the location of each persisted partition's data

Persistence Storage levels

- ❑ Storage location—where is the data stored?
 - MEMORY_ONLY: Store data in memory if it fits
 - DISK_ONLY: Store all partitions on disk
 - MEMORY_AND_DISK: Store any partition that does not fit in memory on disk
 - Called spilling

```
from pyspark import StorageLevel
myDF.persist(StorageLevel.DISK_ONLY)
```

Language: Python

```
import org.apache.spark.storage.StorageLevel
myDF.persist(StorageLevel.DISK_ONLY)
```

Language: Scala

Persistence Storage levels - Memory Serialization

- ❑ In Python, data in memory is always serialized
- ❑ In Scala, you can choose to serialize data in memory
 - By default, in Scala and Java, data in memory is stored objects
 - Use `MEMORY_ONLY_SER` and `MEMORY_AND_DISK_SER` to serialize the objects into a sequence of bytes instead
 - Much more space efficient but less time efficient
- ❑ Datasets are serialized by Spark SQL encoders, which are very efficient
 - Plain RDDs use native Java/Scala serialization by default
 - Use Kryo instead for better performance
- ❑ Serialization options do not apply to disk persistence
 - Files are always in serialized form by definition

Persistence Storage levels - Default Storage Levels

- ❑ The storageLevel parameter for the DataFrame, Dataset, or RDD persist operation is optional
 - The default for DataFrames and Datasets is MEMORY_AND_DISK
 - The default for RDDs is MEMORY_ONLY
- ❑ persist with no storage level specified is a synonym for cache
 - myDF.persist() is equivalent to myDF.cache()

Table and view storage level is always MEMORY_ONLY

Persistence Storage levels - When and Where to Persist

- ❑ When should you persist a DataFrame, Dataset, or RDD?
 - When the data is likely to be reused
 - Such as in iterative algorithms and machine learning
 - When it would be very expensive to recreate the data if a job or node fails

Persistence Storage levels - When and Where to Persist

- ❑ How to choose a storage level
 - Memory—use when possible for best performance
 - Save space by serializing the data if necessary
 - Disk—use when re-executing the query is more expensive than disk read
 - Such as expensive functions or filtering large datasets
 - Replication—use when re-execution is more expensive than bandwidth

Contents

1. **Course Introduction**
2. **Why Apache Spark?**
3. **Spark Cluster Managers**
4. **Introduction to DataBricks**
5. **DataBricks Components**
6. **RDD**
7. **Transformations and Actions in RDD**
8. **DataFrame**
9. **Transformation and Actions in Dataframe**
10. **Working with DataFrames**
11. **SparkSQL**
12. **File systems and sources supported by Spark**
13. **DeltaLake**
14. **Spark Applications**
15. **Batch ETL using Spark**
16. **Introduction to Kafka**
17. **Real-Time ETL and Event partition using Kafka and Spark**
18. **Spark MLLib and Machine Learning using Spark**

Spark Streaming

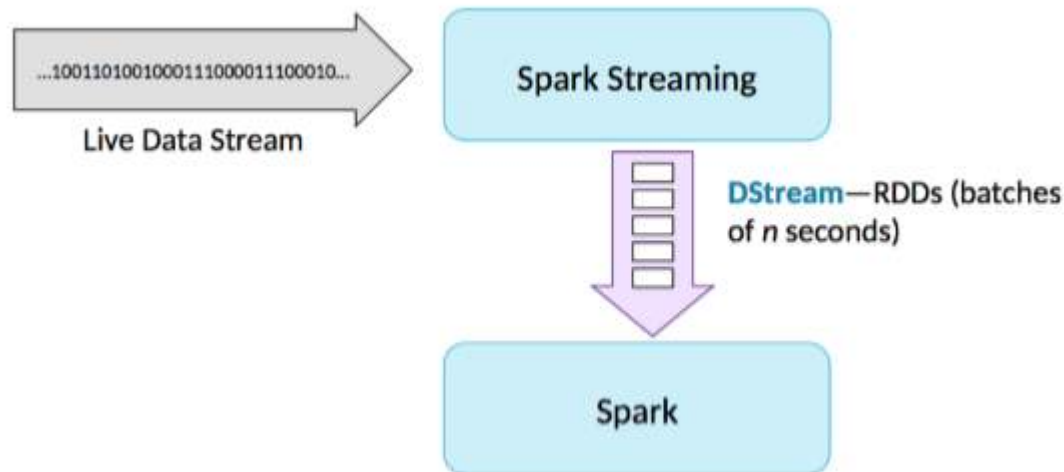
- ❑ An extension of core Spark
- ❑ Provides real-time data processing
- ❑ Segments an incoming stream of data into micro-batches
- ❑ Many big data applications need to process large data streams in real time, such as
 - Continuous ETL
 - Website monitoring
 - Fraud detection
 - Advertisement monetization
 - Social media analysis
 - Financial market trends

Spark Streaming - Features

- ❑ Latencies of a few seconds or less
- ❑ Scalability and efficient fault tolerance
- ❑ “Once and only once” processing
- ❑ Integrates batch and real-time processing
- ❑ Uses the core Spark RDD API

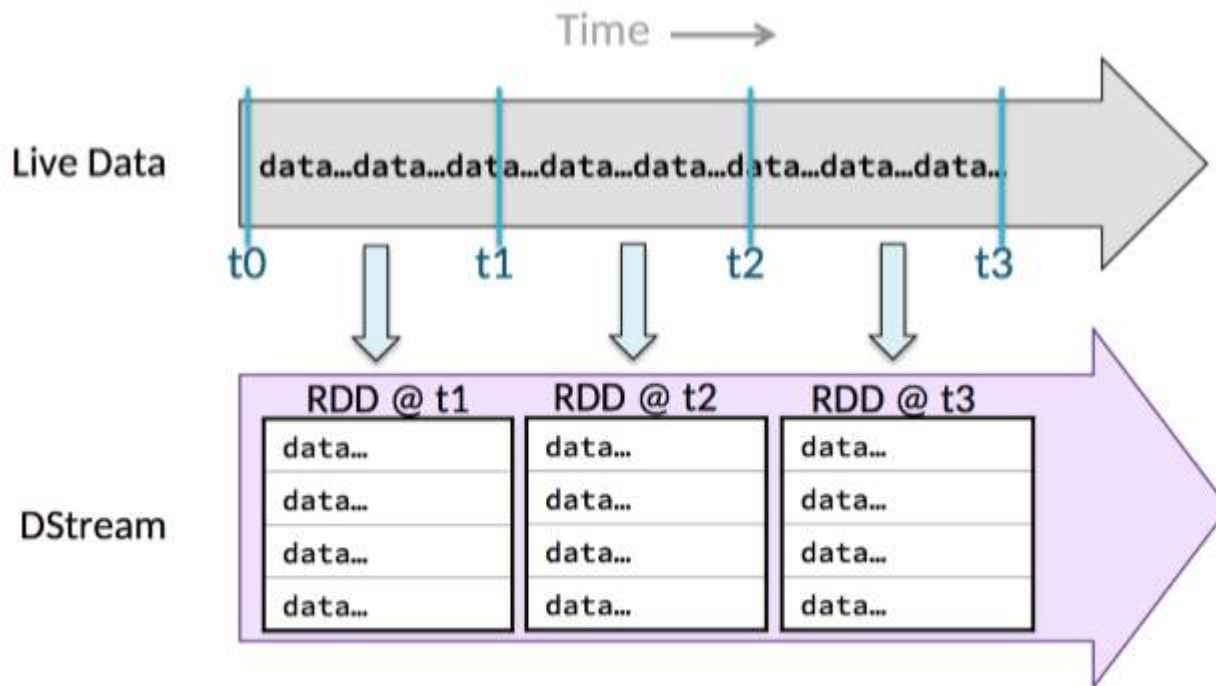
DStreams

- ❑ DStream—RDDs (batches of n seconds)
- ❑ A DStream (Discretized Stream) divides a data stream into batches of n seconds
- ❑ Processes each batch in Spark as an RDD
- ❑ Returns results of RDD operations in batches



DStreams

- ❑ A DStream is a sequence of RDDs representing a data stream



DStreams Data Sources

- ❑ DStreams are defined for a given input stream (such as a Unix socket) — Created by the Streaming context

`ssc.socketTextStream(hostname, port)`

- Similar to how RDDs are created by the Spark context

- ❑ Out-of-the-box data sources

- Network

- Sockets

- Services such as Apache Flume, Apache Kafka, ZeroMQ, or Amazon Kinesis

- Files

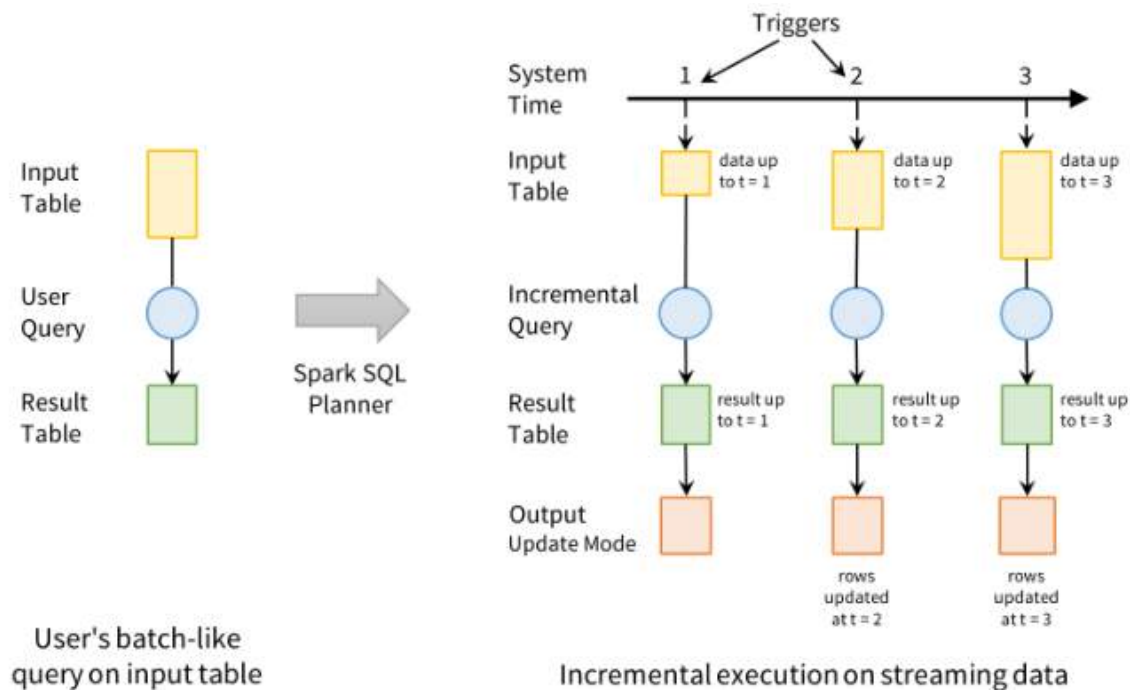
- Monitors an HDFS directory for new content

DStreams Data Sources

- ❑ DStream operations are applied to every RDD in the stream
 - Executed once per duration
- ❑ Two types of DStream operations
 - Transformations
 - Create a new DStream from an existing one
 - Output operations
 - Write data (for example, to a file system, database, or console)
 - Similar to RDD actions

All the other operations are the same as RDD.

Spark Structured Streaming



Spark Structured Streaming

Spark Structured Streaming (aka Structured Streaming or Spark Streams) is the module of Apache Spark for stream processing using streaming queries

Streaming queries can be expressed using a high-level declarative streaming API or SQL (SQL over stream / streaming SQL). The declarative streaming and SQL are executed on the underlying highly-optimized Spark SQL engine.

Spark Structured Streaming

Spark Structured Streaming comes with two stream execution engines for executing streaming queries

- Micro Batch Execution

- Continuous Streaming

Spark Structured Streaming

Using Structured streaming we can posses

- Streaming Aggregation

- Streaming Join

- Streaming Watermark

- Stateful Stream processing

Spark Structured Streaming

In Structured Streaming, Spark developers describe custom streaming computations in the same way as with Spark SQL.

Internally, Structured Streaming applies the user-defined structured query to the continuously and indefinitely arriving data to analyze real-time streaming data.



THANK YOU